

Notes on Randomized Algorithms

Mahdi Zamani
Visa Research, Palo Alto, CA
mzamani@visa.com

1 Basics

Poisson Trials. A sequence of *independent* 0-1 random variables (X_1, \dots, X_n) , where for every $i \in [n]$,

$$\begin{aligned}\Pr(X_i \text{ succeeds}) &= \Pr(X_i = 1) = p_i \\ \Pr(X_i \text{ fails}) &= \Pr(X_i = 0) = 1 - p_i \\ 0 &\leq p_i \leq 1.\end{aligned}$$

Bernoulli Trials. A special case of Poisson trials, where the independent random variables have the same distribution, i.e., for every $i \in [n]$,

$$\begin{aligned}\Pr(X_i = 1) &= p \\ \Pr(X_i = 0) &= 1 - p \\ 0 &\leq p \leq 1.\end{aligned}$$

Binomial Distribution. The distribution of the number of successes in a sequence of Bernoulli trials, e.g., the number of heads in a sequence of n coin flips. A binomial random variable X with parameters n and p , denoted by $B(n, p)$, is defined by the following probability distribution on $j = 0, 1, \dots, n$:

$$\begin{aligned}\Pr(X = j) &= \binom{n}{j} p^j (1 - p)^{n-j} \\ \mathbf{E}[X] &= np.\end{aligned}$$

Geometric Distribution. The distribution of the number of trials in a Bernoulli process until we get a success, e.g., the number of times to flip a coin until it lands on heads. A geometric random variable X with parameter p is defined by the following probability distribution:

$$\begin{aligned}\Pr(X = n) &= p(1 - p)^{n-1} \\ \mathbf{E}[X] &= 1/p.\end{aligned}$$

Poisson Distribution. A Poisson random variable X with parameter μ is defined by the following probability distribution on $j = 0, 1, \dots, n$:

$$\Pr(X = j) = \frac{e^{-\mu} \mu^j}{j!}.$$

- The Poisson distribution expresses the probability of a given number of events occurring in a fixed interval of time or space if these events occur with a known constant rate and independently of the time since the last event.
- When throwing m balls into n bins, the distribution of the number of balls in a bin is approximately Poisson with $\mu = m/n$ which is exactly the average number of balls per bin.

Exponential Distribution. The exponential distribution is *memoryless* because the past has no impact on its future behavior.

Bounds on Poisson Random Variables. Let X be a Poisson random variable with parameter μ .

1. If $x > \mu$, then

$$\Pr(X \geq x) \leq \frac{e^{-\mu}(e\mu)^x}{x^x}.$$

2. If $x < \mu$, then

$$\Pr(X \leq x) \leq \frac{e^{-\mu}(e\mu)^x}{x^x}.$$

Properties of Expectations. • If X is a random variable and a is a constant, then $\mathbf{E}[aX] = a\mathbf{E}[X]$. Thus, $\mathbf{E}[\mathbf{E}[X]] = \mathbf{E}[X]$, because $\mathbf{E}[X]$ is a constant.

- Linearity of expectation: For any set of random variables X_1, \dots, X_n , we have $\mathbf{E}[\sum_{i=1}^n X_i] = \sum_{i=1}^n \mathbf{E}[X_i]$.

2 Concentration Bounds

Union Bound (Boole's Inequality). The probability that at least one event in a set of events happens is no greater than the sum of the probabilities of the individual events, i.e., for a set of events A_1, \dots, A_n ,

$$\Pr\left(\bigcup_{i=1}^n A_i\right) \leq \sum_{i=1}^n \Pr(A_i).$$

Markov's Inequality. Let X be a random variable that assumes only non-negative values. Then, for all $a > 0$,

$$\Pr(X \geq a) \leq \frac{\mathbf{E}[X]}{a}.$$

Variance. The variance of a random variable X is defined as $\mathbf{Var}[X] = \mathbf{E}[X^2] - (\mathbf{E}[X])^2$.

If X is a binomial random variable, then $\mathbf{Var}[X] = np(1-p)$.

If X is a geometric random variable, then $\mathbf{Var}[X] = (1-p)/p^2$.

Chebyshev's Inequality. For any $a > 0$,

$$\Pr(|X - \mathbf{E}[X]| \geq a) \leq \frac{\mathbf{Var}[X]}{a^2}.$$

Chernoff Bounds. Let X_1, \dots, X_n be independent Poisson trials such that $\Pr(X_i) = p_i$. Let $X = \sum_{i=1}^n X_i$, $\mu = \mathbf{E}[X]$, and $\mu_L \leq \mu \leq \mu_H$. The following Chernoff bounds hold:

$$\text{For } \delta > 0, \quad \Pr(X \geq (1 + \delta)\mu) \leq \left(\frac{e^\delta}{(1+\delta)^{(1+\delta)}} \right)^\mu \quad [?, \text{ Page 64}]$$

$$\text{For } 0 < \delta < 1, \quad \Pr(X \leq (1 - \delta)\mu) \leq \left(\frac{e^{-\delta}}{(1-\delta)^{(1-\delta)}} \right)^\mu \quad [?, \text{ Page 66}]$$

$$\text{For } 0 < \delta < 1, \quad \Pr(X \geq (1 + \delta)\mu) \leq e^{-\mu\delta^2/3} \quad [?, \text{ Page 64}]$$

$$\text{For } 0 < \delta < 1, \quad \Pr(X \leq (1 - \delta)\mu) \leq e^{-\mu\delta^2/2} \quad [?, \text{ Page 66}]$$

$$\text{For } 0 < \delta < 1, \quad \Pr(|X - \mu| \geq \delta\mu) \leq 2e^{-\mu\delta^2/3} \quad [?, \text{ Page 67}]$$

$$\text{For } \delta > 0, \quad \Pr(X < \mu - \delta), \Pr(X > \mu + \delta) \leq e^{-2\delta^2/n} \quad [?, \text{ Page 6}]$$

$$\text{For } \delta > 0, \quad \Pr(X < \mu_L - \delta), \Pr(X > \mu_H + \delta) \leq e^{-2\delta^2/n} \quad [?, \text{ Page 8}]$$

$$\text{For } \delta > 2e\mu, \quad \Pr(X > \delta) \leq 2^{-\delta} \quad [?, \text{ Page 7}]$$

$$\text{For } \delta > 1, \quad \Pr(X \geq (1 + \delta)\mu) \leq e^{-\mu\delta/3} \quad \text{Wikipedia}$$

Convex Function. A function $f : \mathbb{R} \rightarrow \mathbb{R}$ is said to be convex if and only if, for any x_1, x_2 and $0 \leq \lambda \leq 1$,

$$f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2).$$

Jensen's Inequality. If f is a convex function, then $\mathbf{E}[f(X)] \geq f(\mathbf{E}[X])$.

Conditional Expectation. A conditional expectation is not a scalar value; it is a random variable defined as

$$\mathbf{E}[X|Y] = \{\mathbf{E}[X|Y = y] \text{ with } \Pr(Y = y)\}.$$

In other words, $\mathbf{E}[X|Y]$ is a function that maps every $y \in Y$ to a $x \in X$. It is easy to see that

$$\mathbf{E}[\mathbf{E}[X|Y]] = \mathbf{E}[X].$$

Also,

$$\mathbf{E}[\mathbf{E}[X|Y, Z]|Z] = \mathbf{E}[X|Z].$$

3 Martingales

A martingale is a sequence of random variables used to model an event that is a function of a sequence of past events. Martingales are useful for events where the knowledge of past events does not allow to predict the expected value or the actual probability of the event.

For example, suppose that we are throwing m balls independently and uniformly at random into n bins. Let X_i be the random variable representing the bin into which the i -th ball falls, and F be the number of empty bins after the m balls are thrown. Then, the sequence

$$Z_i = \mathbf{E}[F | X_0, \dots, X_i]$$

is a martingale. This model allows us to find a concentration bound for F using Azuma's inequality (which we will describe later):

$$\Pr(|F - \mathbf{E}[F]| \geq \epsilon) \leq 2e^{-2\epsilon^2/m}.$$

Definition 1. A sequence of random variables Z_0, Z_1, \dots is a martingale with respect to a sequence X_0, X_1, \dots , if for all $i \geq 0$, the following conditions hold:

1. Z_i is a function of X_0, \dots, X_i ;
2. $\mathbf{E}[|Z_i|] < \infty$; and
3. $\mathbf{E}[Z_{i+1} | X_0, \dots, X_i] = Z_i$.

Definition 2. A sequence of random variables Z_0, Z_1, \dots is called a martingale when it is a martingale with respect to itself, that is

1. $\mathbf{E}[|Z_i|] < \infty$
2. $\mathbf{E}[Z_{i+1} | Z_0, \dots, Z_i] = Z_i$

Definition 3. Let X_0, \dots, X_n be a sequence of random variables and Y be a random variable that depends on the X_i 's. Then, for $i = 0, 1, \dots, n$,

$$Z_i = \mathbf{E}[Y | X_0, \dots, X_i]$$

is a martingale with respect to X_0, \dots, X_n and is called a Doob martingale.

Definition 4. A random variable T is a stopping time for $\{Z_i\}$ if the event $T = n$ depends only on Z_0, \dots, Z_n .

Theorem 1. [Martingale Stopping Time Theorem] If Z_0, Z_1, \dots is a martingale with respect to X_0, X_1, \dots and T is a stopping time for X_1, X_2, \dots , then

$$\mathbf{E}[Z_T] = \mathbf{E}[Z_0],$$

if one of the following holds:

1. $|Z_i| < c$ for some constant c .
2. T is bounded.

3. $\mathbf{E}[T] < \infty$ and there exists a c such that $E[Z_{i+1} - Z_i | X_1, \dots, X_i] < c$.

Theorem 2. [Azuma's Inequality] *Let X_0, \dots, X_n be a martingale such that $|X_k - X_{k-1}| \leq c_k$. Then, for all $t \geq 0$ and any $\lambda > 0$,*

$$\Pr(|X_t - X_0| \geq \lambda) \leq 2e^{-\lambda^2 / (2 \sum_{k=1}^t c_k^2)}.$$

Theorem 3. [Azuma's Inequality – Tighter Bound] *Let X_0, \dots, X_n be a martingale such that*

$$B_k \leq X_k - X_{k-1} \leq B_k + d_k$$

for some constants d_k and for some random variables B_k . Then, for all $t \geq 0$ and any $\lambda > 0$,

$$\Pr(|X_t - X_0| \geq \lambda) \leq 2e^{-2\lambda^2 / (\sum_{k=1}^t d_k^2)}.$$

Definition 5. [Informal] *A dependency graph for a set of events E_1, \dots, E_n is graph with a vertex for each event, where there is an edge between two vertices iff they have some dependency with each other. In other words, if the two vertices are mutually independent, then there is no edge between them.*

Theorem 4. [Lovasz Local Lemma] *Let E_1, \dots, E_n be a set of events and*

1. *for all i , $\Pr(E_i) \leq p$;*
2. *the degree of the dependency graph of E_1, \dots, E_n is at most d ;*
3. *$ep(d+1) \leq 1$.*

Then,

$$\Pr\left(\bigcap_{i=1}^n \bar{E}_i\right) > 0.$$

4 Useful bounds

The following bounds follow from the Taylor series expansion of e^x :

$$1 + x \leq e^x$$

$$1 - x \leq e^{-x}$$

$$\text{For } x \geq 0 : 1 + x \geq e^x(1 - x^2)$$

$$\text{For } x \geq 0 : 1 - x \geq e^{-x} - x^2/2$$

References